

Pivotal.

SpringOne Platform

by Pivotal.

# re:Cap Seoul 2019

2019년 12월 17일(화) • 한화 드림플러스 강남

# Apache Geode Summit re:Cap

신혜원 (Haewon Shin)  
Pivotal Platform Architect



# Apache Geode Summit

---



## Apache Geode Summit 2019

<https://www.youtube.com/playlist?list=PLgCUiLdOC2pi9SYptnVamh6xxoeRVZx8s>

# Geode Sessions


- **Breaking Open Apache Geode: How It Works and Why**
- Introducing the Geode Native Client
- **Performance in Geode: How Fast Is It, How Is It Measured, and How Can It Be Improved?**
- Using Apache Geode: Lessons Learned at Southwest Airlines
- A Fireside Chat with Apache Geode Committers
- **Visualize Your Geode Metrics**
- **Reactive Event Processing with Apache Geode**
- Data Serialization and CI/CD Techniques for Apache Geode
- High-Performance Data Processing with Spring Cloud Data Flow and Geode
- Scaling Beyond a Billion Transactions Per Day with Sub-Second Responses
- Scalable, Cloud-Native Data Applications by Example
- Simple Data Movement Patterns: Legacy Application to Cloud-Native Environment and Apache Geode

# Agenda

---

- What is Geode?
- How to Monitor: Visualize Geode Metrics
- Performance in Geode
- Reactive Programming with Geode

# What is Geode?



# 피보탈이 투자하고 있는 오픈소스 생태계



# Geode / GemFire / Pivotal Cloud Cache (PCC)

Continuous investment in data R&D

What is good for one is good for the other



**Pivotal  
Cloud Cache**

On-platform & clouds

High-performance caching & data  
acceleration for microservices

<https://pivotal.io/pivotal-cloud-cache>



**Pivotal  
GemFire®**

Off-platform & native clients

In-memory data grid enabling  
event-driven architecture and  
fast-data access patterns

<https://pivotal.io/pivotal-gemfire>

Open source - Performance is key. Consistency is a must.


Reliable transaction processing and shared-nothing architecture for very low  
latency performance with high concurrency processing

<https://content.pivotal.io/pivotal-gemfire/scaling-data-services-with-pivotal-gemfire>




# What is Geode?

- Distributed key-value store




# What is Geode?

- Distributed key-value store
- High available



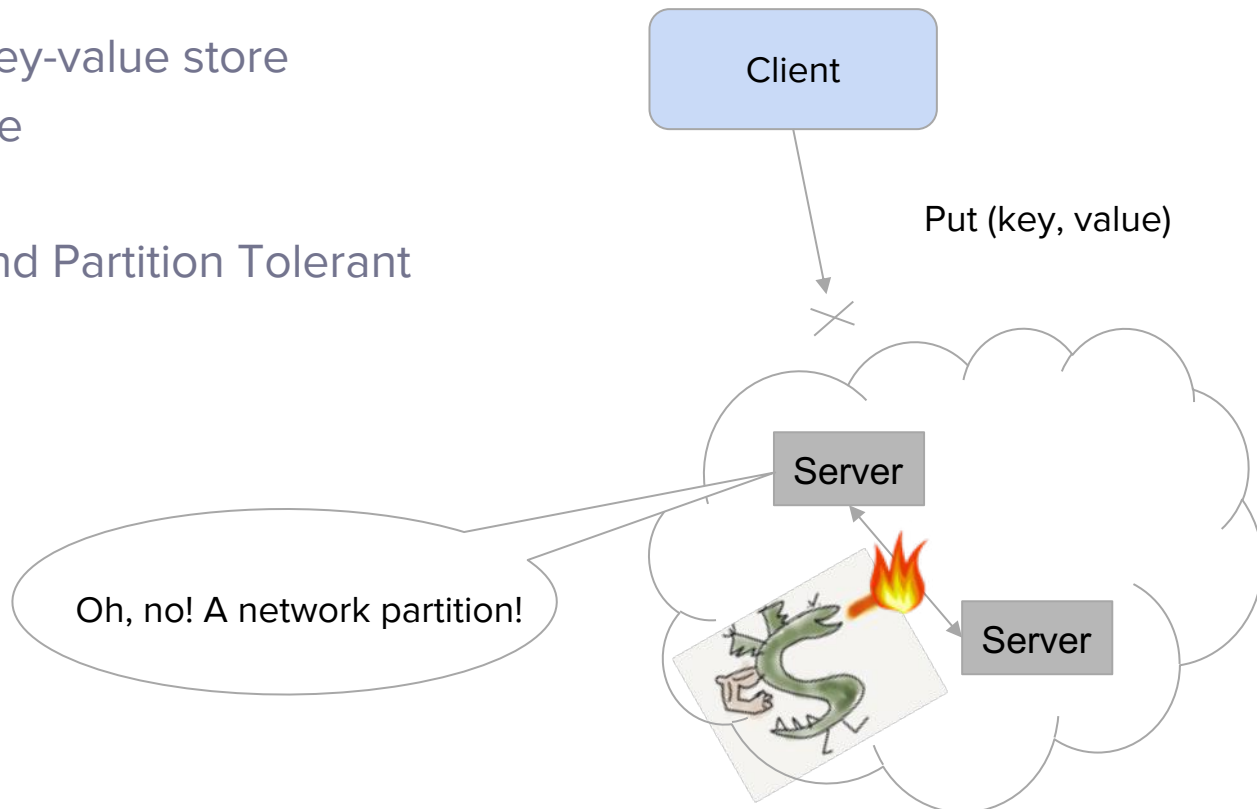
# What is Geode?

- Distributed key-value store
- High available
- Low Latency



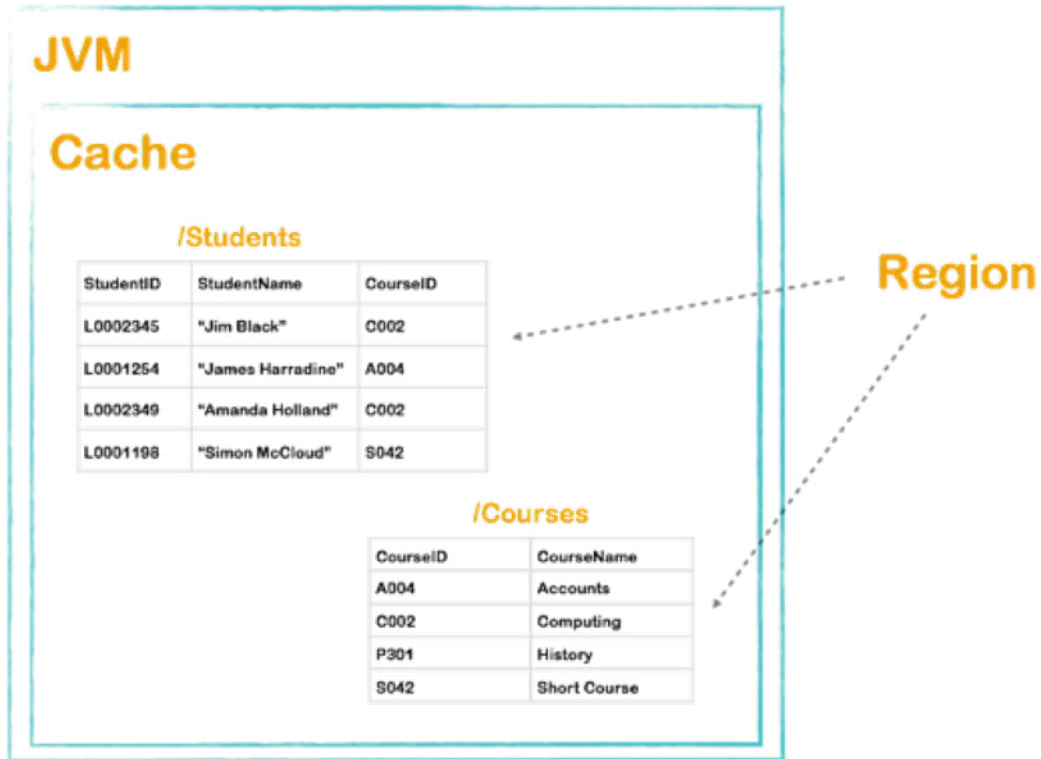
# What is Geode?

- Distributed key-value store
- High available
- Low Latency
- Consistent and Partition Tolerant




# Regions

- Synonymous to a Table in NoSQL terminology
- Stores Data in <Key,Value> pairs with unique Keys
- Divided into buckets across Cache Members




# Regions Data Placement Design

## Replicated



## Partitioned



# Visualize your Geode Metrics



# Viewing Geode Metrics Before Now

---

- Internal metrics written to a local file in a proprietary format
  - Viewable with a custom viewing tool
  - Generally viewed after the events happened
- A subset published via JMX



192.168.0.60

Cluster View

Data Browser

All

**Normal**  
Cluster Status17.78 GB  
Total Heap

4

Members

3

Servers

0

Clients

1

Locators

2

Regions

0

Functions

0


Unique CQs

0

Subscriptions

Members

Data

☒ Topology☐ Server Groups☐ Redundancy Zones

Managers



Locators



Servers

MEMORY USAGE

413.96 MB

Last 15 Minutes



DISK THROUGHPUT

Last 15 Minutes

Reads  
Writes

0.00

Reads

0.00

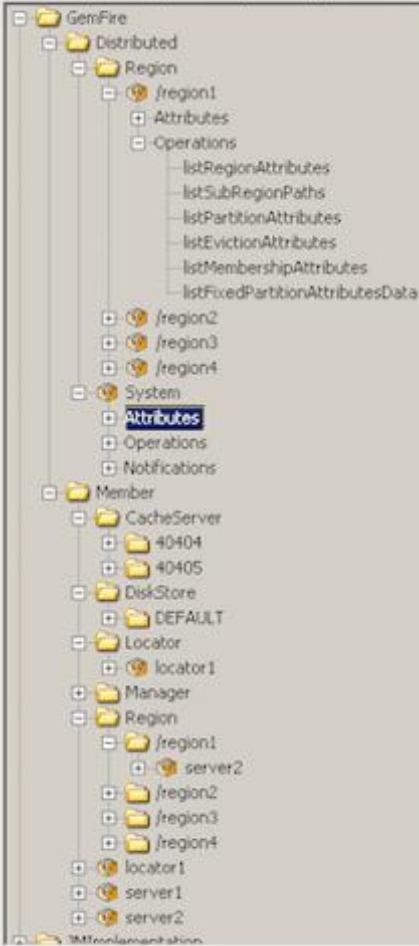
Writes

KEY STATISTICS



NO. OF JVM PAUSES





## Attribute values

| Name                       | Value   |
|----------------------------|---|
| ActiveCQCount              | 0   |
| AlertLevel                 | severe  |
| AverageReads               | 0.0   |
| AverageWrites              | 0.0   |
| DiskFlushAvgLatency        | 0   |
| DiskReadsRate              | 0.0   |
| DiskWritesRate             | 0.0   |
| DistributedSystemId        | -1  |
| GarbageCollectionCount     | 1509428   |
| JVMPauses                  | 6   |
| LocatorCount               | 1   |
| ManagerObjectName          | GemFire:service=Manager,type=Member,member=locator1 |
| MemberCount                | 3   |
| MemberObjectName           | GemFire:type=Member,member=locator1                 |
| NumClients                 | 0   |
| NumInitialImagesInProgress | 0   |
| NumRunningFunctions        | 0   |
| NumSubscriptions           | 0   |
| QueryRequestRate           | 0.0   |
| RegisteredCQCount          | 0   |
| SystemDiskStoreCount       | 2   |
| TotalBackupInProgress      | 0   |
| TotalBytesOnDisk           | 0   |
| TotalDiskUsage             | 0   |
| TotalHeapSize              | 369   |
| TotalHitCount              | 35  |
| TotalMissCount             | 7   |
| TotalRegionCount           | 4   |
| TotalRegionEntryCount      | 6   |
| UsedHeapSize               | 19  |

Refresh

File: 1. 12\_Delta\_Solution/stats1007 Windows 7 6.1 x86 blr-bchoudhury GemFire 6.6.2 #build 34761 as of 03/06/2012 17:10:10 PST

| StartTime      | File | Samples | Pid  | Type                 | Name   |
|----------------|------|---------|------|----------------------|--|
| 06/08 16:09:56 | 1    | 78      | 7468 | vmGCStats            | Copy   |
| 06/08 16:09:56 | 1    | 78      | 7468 | StatSampler          | statSampler  |
| 06/08 16:09:56 | 1    | 77      | 7468 | ResourceManagerStats | ResourceManagerStats                                   |
| 06/08 16:09:56 | 1    | 77      | 7468 | PoolStats            | client->[anyServers]                                   |
| 06/08 16:09:56 | 1    | 77      | 7468 | ClientStats          | ClientStats-client-blr-bchoudhury.vmware.com:40406     |
| 06/08 16:09:56 | 1    | 78      | 7468 | ClientSendStats      | ClientSendStats-client-blr-bchoudhury.vmware.com:40406 |
| 06/08 16:09:56 | 1    | 77      | 7468 | CachePerfStats       | RegionStats-InventoryItem                              |
| 06/08 16:09:56 | 1    | 77      | 7468 | CachePerfStats       | RegionStats-Customer                                   |
| 06/08 16:09:56 | 1    | 77      | 7468 | Cache                |  |
| 06/08 16:09:56 | 1    | 77      | 7468 | Cache                |  |
| 06/08 16:09:56 | 1    | 77      | 7468 | Cache                |  |
| 06/08 16:09:56 | 1    | 77      | 7468 | Cache                |  |
| 06/08 16:09:56 | 1    | 77      | 7468 | Cache                |  |

## cacheListenerCallsCompleted

cacheListenerCallsInProgress

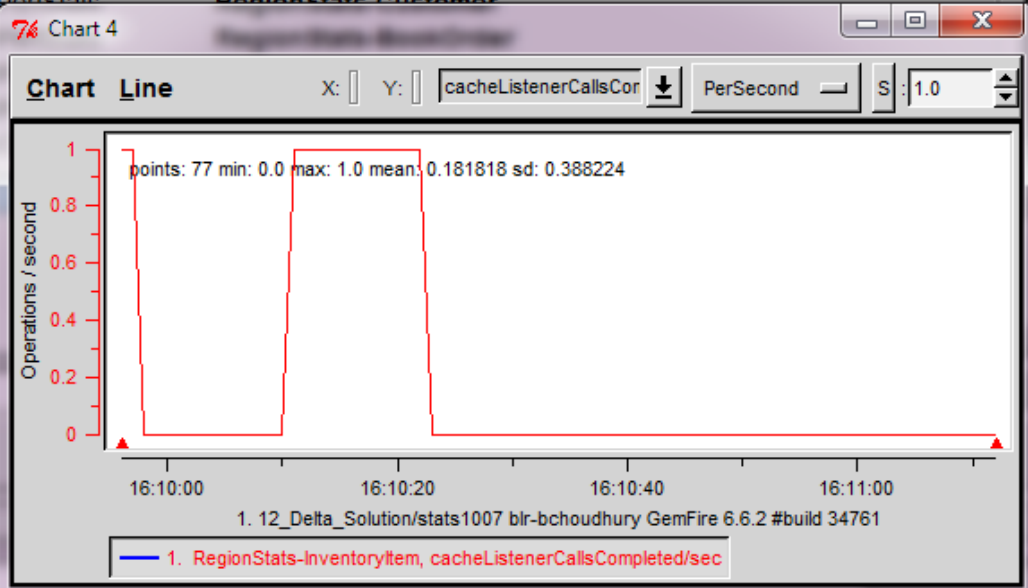
cacheListenerCallTime

cacheWriterCallsCompleted

cacheWriterCallsInProgress

Chart: Chart 4

Add Line

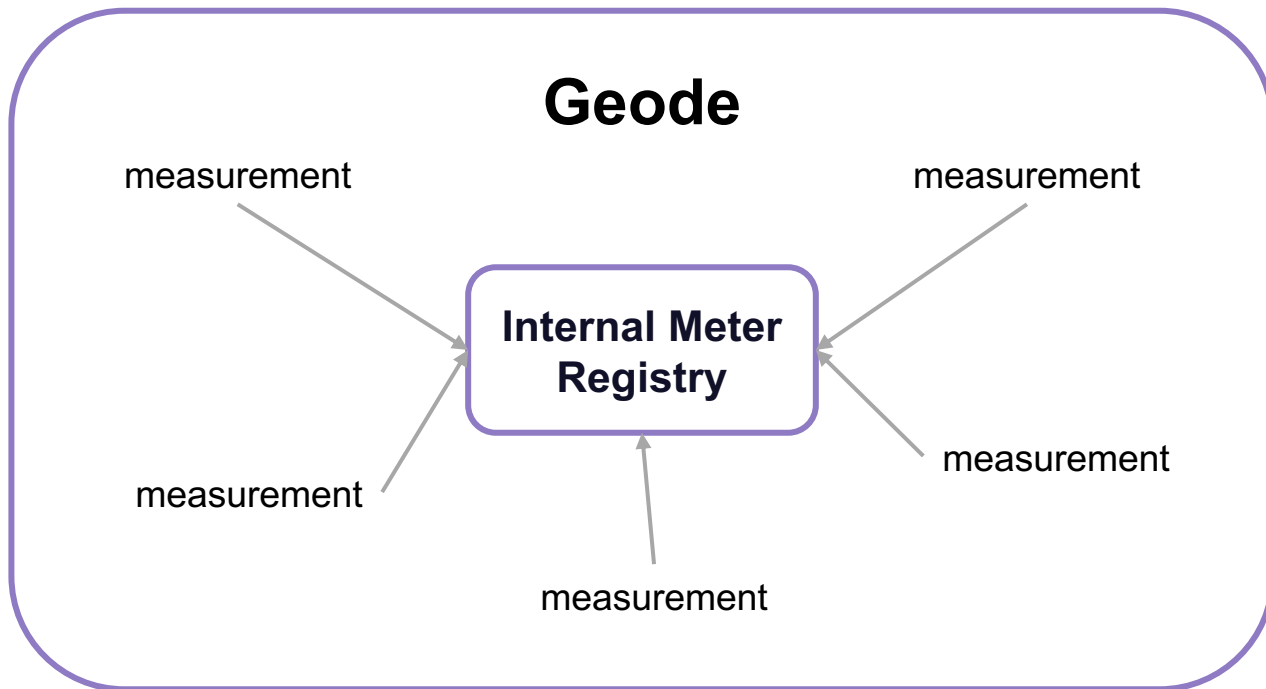


# Micrometer in Geode


---

- Make key metrics visible in external monitoring systems
- View metrics while the system is running
- Augment metrics with details to aid understanding

# Meter Registries



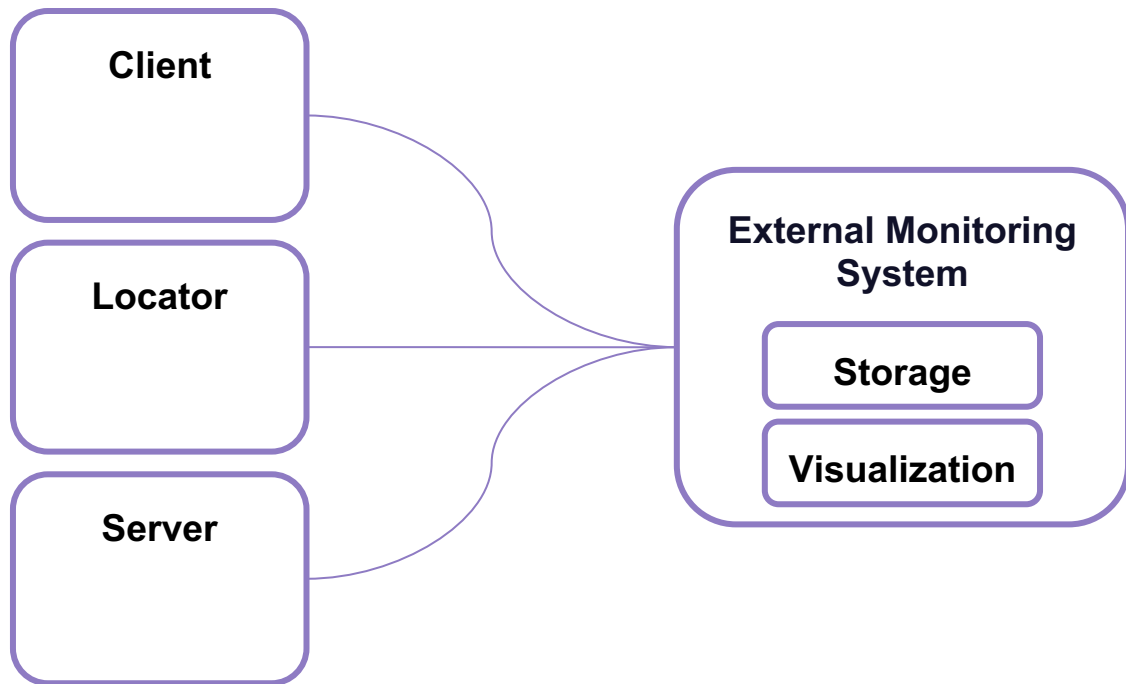
# Connecting to Monitoring Systems



# Connecting to Monitoring Systems

- Chose a Micrometer meter registry implementation that publishes to your monitoring system
  - For a list of implementations, see:  
<http://micrometer.io/docs>
- Create a MetricsPublishingService that adds your registry to Geode

# Live View of Metrics








Cache Entries



Get Max Latency




CPU Usage



Put Max Latency



Memory Usage



# How Geode Records Measurements


- Uses Micrometer *meters* to record measurements
  - `eventsReceivedCounter.increment(delta);`
- Collects the meters in its *meter registry*

# Current Geode Meters


---

- System metrics
  - Processors
  - File descriptors
- JVM metrics
  - Process
  - Memory
  - Garbage collection
  - Threads
- Cache Entries


## jvm.memory.used





# Scenario – Partition Region Redundancy 1



geode.cache.entries  
{region="my-partitioned-region"}








Cache Entries



Get Max Latency



CPU Usage



Put Max Latency



Memory Usage





# The Future Depends on You

---

- Send Geode's metrics to your favorite monitoring system
- What works well for you?
- What additional metrics would best help you manage your Geode systems?

# For Further Information

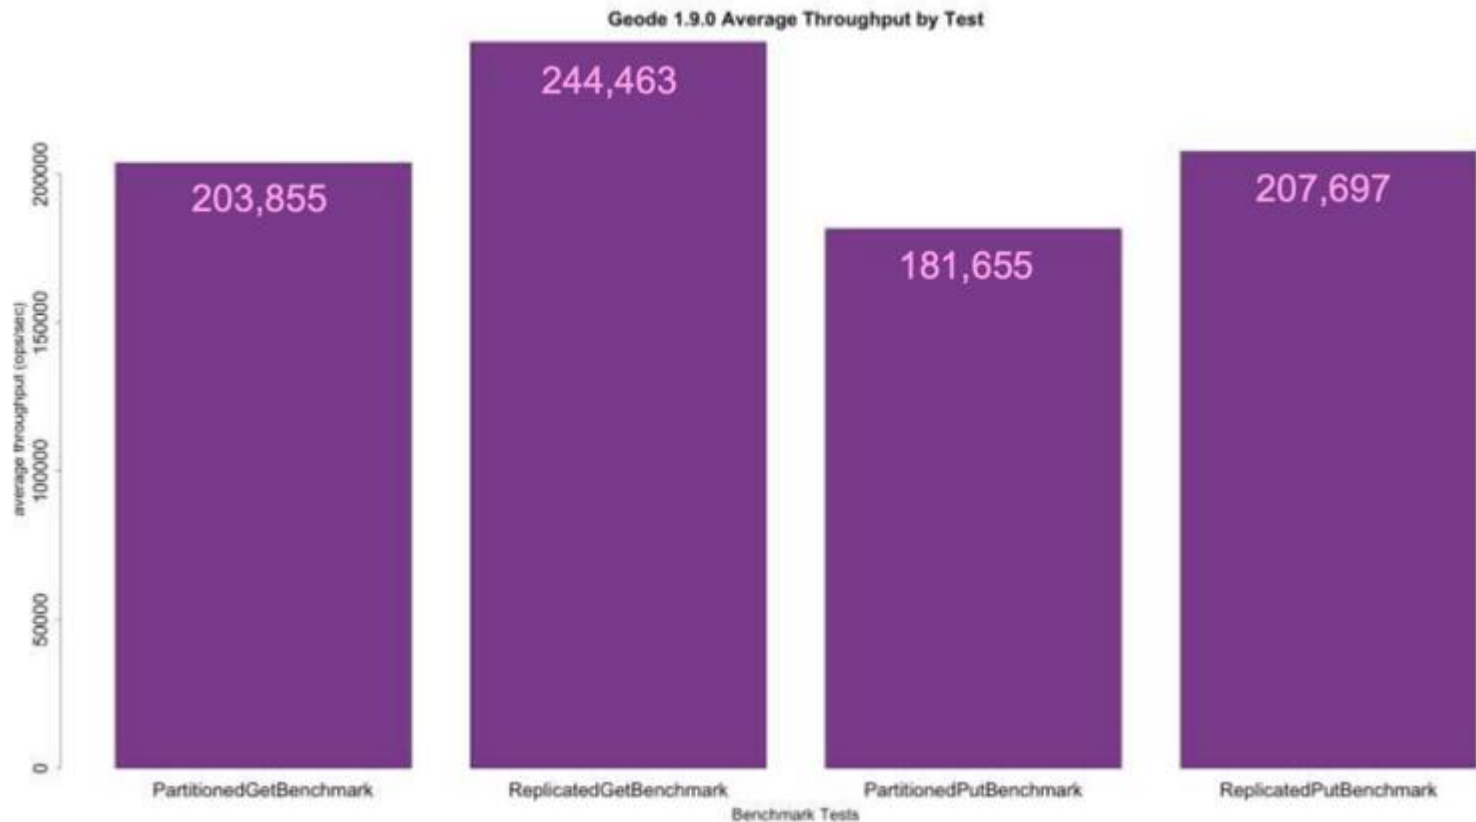
---

- Micrometer: <http://micrometer.io/>
- Prometheus: <https://prometheus.io>
- Grafana: <https://grafana.com>
- Publishing Geode Metrics to External Monitoring Systems:  
<https://cwiki.apache.org/confluence/display/GEODE/Publishing+Geode+Metrics+to+External+Monitoring+Systems>
- Example code to publish metrics to Prometheus:  
<https://github.com/moleske/geode-registry-example>

# Performance in Geode



# Performance of Geode 1.9.0



# Creating the Geode Benchmark – Features

---

- Run by anyone interested in Geode
- Have others create benchmarks
- Visualize benchmark results over time
- Increase benchmark coverage of Geode

# Creating the Geode Benchmark – Goals

---

- On demand
- Against any revision of Geode
- On AWS cluster deployment of Geode
- On any dev machine in the office
- From Concourse CI pipeline
- With a profiler attached
- Compare two runs of benchmarks for performance changes

# Tests Currently in the Benchmarks

- **ReplicatedGetBenchmark**
- ReplicatedGetLongBenchmark
- **ReplicatedPutBenchmark**
- ReplicatedPutLongtBenchmark
- ReplicatedPutAllBenchmark
- ReplicatedPutAllLongBenchmark
- ReplicatedFunctionExecutionBenchmark
- ReplicatedFunctionExecutionWithArgumentsBenchmark
- ReplicatedFunctionExecutionWithFiltersBenchmark
- **PartitionedGetBenchmark**
- PartitionedGetLongBenchmark
- **PartitionedPutBenchmark**
- PartitionedPutLongBenchmark
- PartitionedPutAllBenchmark
- PartitionedPutAllLongBenchmark
- PartitionedIndexedQueryBenchmark
- PartitionedFunctionExecutionWithArgumentsBenchmark
- PartitionedFunctionExecutionWithFiltersBenchmark

# Tests Currently in the Benchmarks

---

<https://github.com/apache/geode-benchmarks>



# Other Tested Configurations

---

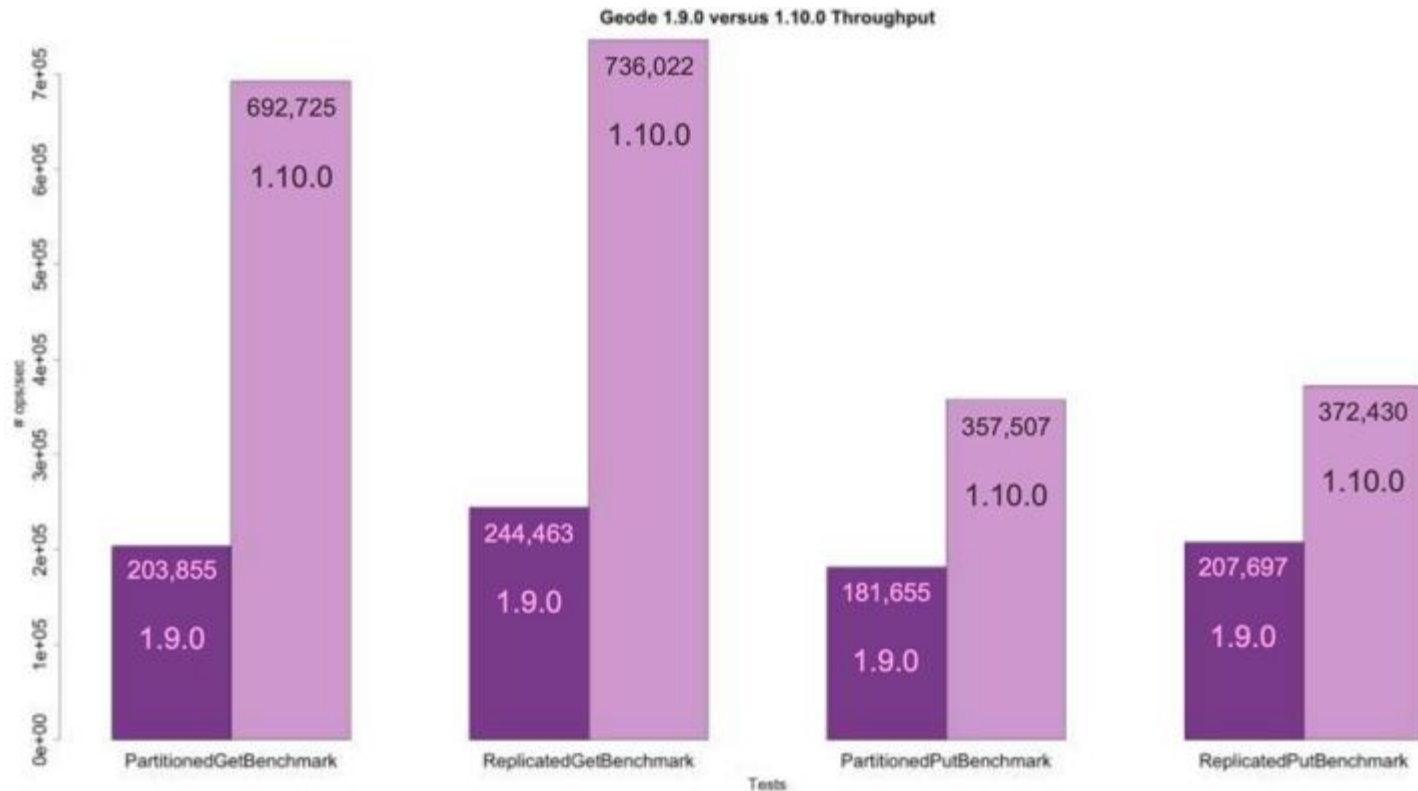
- With SSL
- With JDKs: 8, 11, 12, 13
- With Security Manager
- With Garbage Collectors:
  - CMS
  - G1
  - Z
  - Shenandoah
- Adjustable max heap size

# Finding Performance Bottleneck


---

- Monitor locks
- Thread Park/Unpark Reentrant Locks
- Allocations/GC
- Overuse of synchronization
- Getting a system property in a hot path
- Lazy initialization of objects in a hot path
- Synchronization on a container (ex. Hash map)

# Comparing Performance of 1.9.0 & 1.10.0



# Comparing Performance of 1.9.0 & 1.10.0



# Reactive Event Processing with Geode



# reactive programming is about...

---

“non-blocking, event-driven applications that scale with a small number of threads with backpressure as a key ingredient that aims to ensure producers do not overwhelm consumers”

- Rossen Stoyanchev  
[spring.io](https://spring.io/blog/2016/07/28/reactive-streams-overview) blog July 28, 2016

# Questions

---

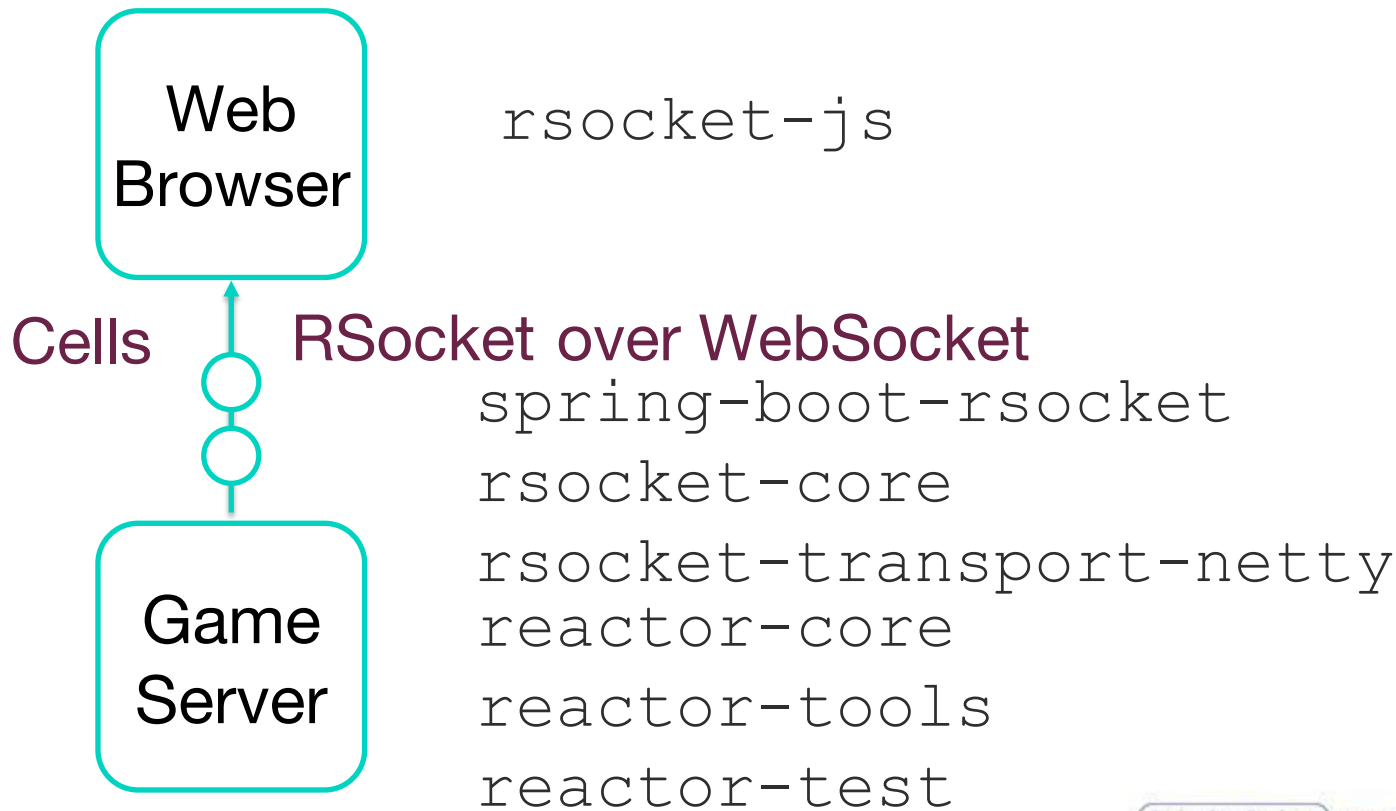
- Geode as reactive consumer (Subscriber)
  - How do active APIs fit with Geode's can Geode's load shedding policy be adapted to reactive back-pressure
- Geode as reactive producer (Publisher)
  - Again, how do reactive APIs fit Geode's
  - Can Geode produce long data streams incrementally

# Conway's Game of Life

```
162 if (wasAlive) {
163     if (liveNeighborsCount < 2) {
164         return Cell.createDead(newCoordinates); // underpopulation
165     } else if (liveNeighborsCount > 3) {
166         return Cell.createDead(newCoordinates); // overpopulation
167     } else {
168         return Cell.createAlive(newCoordinates, isNewborn: false); // survival
169     }
170 } else {
171     if (liveNeighborsCount == 3) {
172         return Cell.createAlive(newCoordinates, isNewborn: true); // reproduction
173     } else {
174         return Cell.createDead(newCoordinates); // status quo
175     }
176 }
```




# A Reactive Toy



```
± |master {1} S:2 U:11 ?:2 ×| → ./gradlew :gameserver:bootRun
```



# Record



spring-boot-rsocket  
rsocket-core  
rsocket-transport-netty  
  
reactor-core  
reactor-tools  
reactor-test

... +  
spring-geode-starter  
spring-data-geode-test  
spring-boot-starter-test

# Parallel putAll

```
175 return Flux.from(source) Flux<Cell>
176     .limitRequest(generations * coordinateSystem.size()) Flux<Cell>
177     .buffer(coordinateSystem.size()) Flux<List<Cell>>
178     .parallel(parallelism) ParallelFlux<List<Cell>>
179     // NB: gotta runOn() after parallel() to actually schedule work in parallel!
180     .runOn(Schedulers.parallel()) // uncomment to demonstrate BlockHound
181     // .runOn(Schedulers.elastic()) // uncomment to satisfy BlockHound
182     .doOnNext(
183         bulkCellConsumer) ParallelFlux<List<Cell>>
184     .sequential() Flux<List<Cell>>
185     .doOnTerminate(summarizePerformance(n, starting, firstElementReceived));
```

# {put,putAll} x {serial,parallel}

|          | put | putAll |
|----------|-----|--------|
| serial   | 8   | 23     |
| parallel | 21  | 124    |

thousands of Cells per second  
MacBook Pro; 8-way parallelism  
de-tuned—for *relative* comparison only

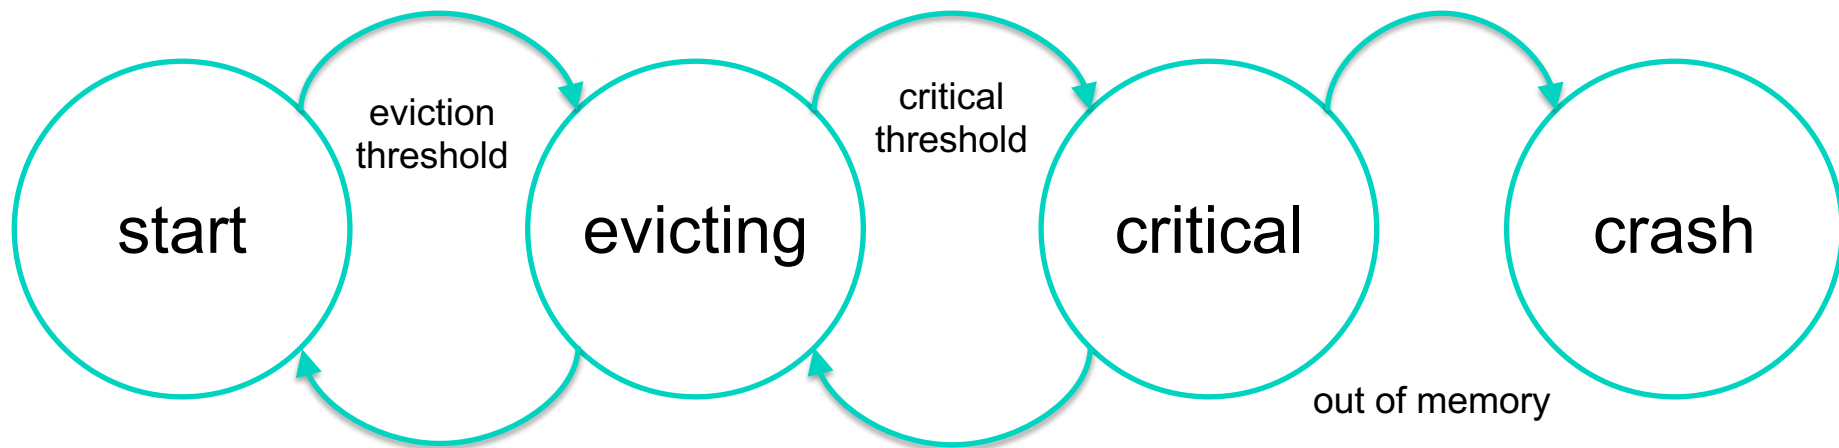
# But remember...

---

“  
reactive programming is about...  
**non-blocking...**”

- Rossen Stoyanchev  
[spring.io](https://spring.io/blog/2016/07/28/reactive-programming-is-about-non-blocking) blog July 28, 2016

# Geode ResourceManager




evictions happening

puts + queries serviced

LOAD SHEDDING

# Eviction Config (purposely fragile)

```
68 factory.setEvictionAttributes(  
69     /*  
70      Region grows until evictionHeapPercentage is reached, then TBD elements are evicted  
71      daemon monitors heap memory usage--non-cache actions can result in eviction  
72      */  
73     EvictionAttributes.createLRUHeapAttributes()  
74 );
```

```
22 @CacheServerApplication(name = "AutoConfiguredContinuousQueryIntegrationTests",  
23     // FRAGILE  
24     // criticalHeapPercentage = 75f, evictionHeapPercentage = 70f)  
25  
26     //  ROBUST-ISH  
27     criticalHeapPercentage = 90f, evictionHeapPercentage = 70f)  
--
```



# Eviction Config (purposely fragile)

```
@BeforeClass
```

```
public static void startGeodeServer() throws IOException {
```

```
    startGemFireServer(
```

```
        // FRAGILE WITHOUT RETRY, ROBUST WITH RETRY
```

```
        GeodeTestServerConfigurationAutoEvictionNoCQ.class,
```

```
        ...arguments: "-Xmx70m", "-Xms70m",
```

```
        // ROBUST (NO RETRY NEEDED)
```

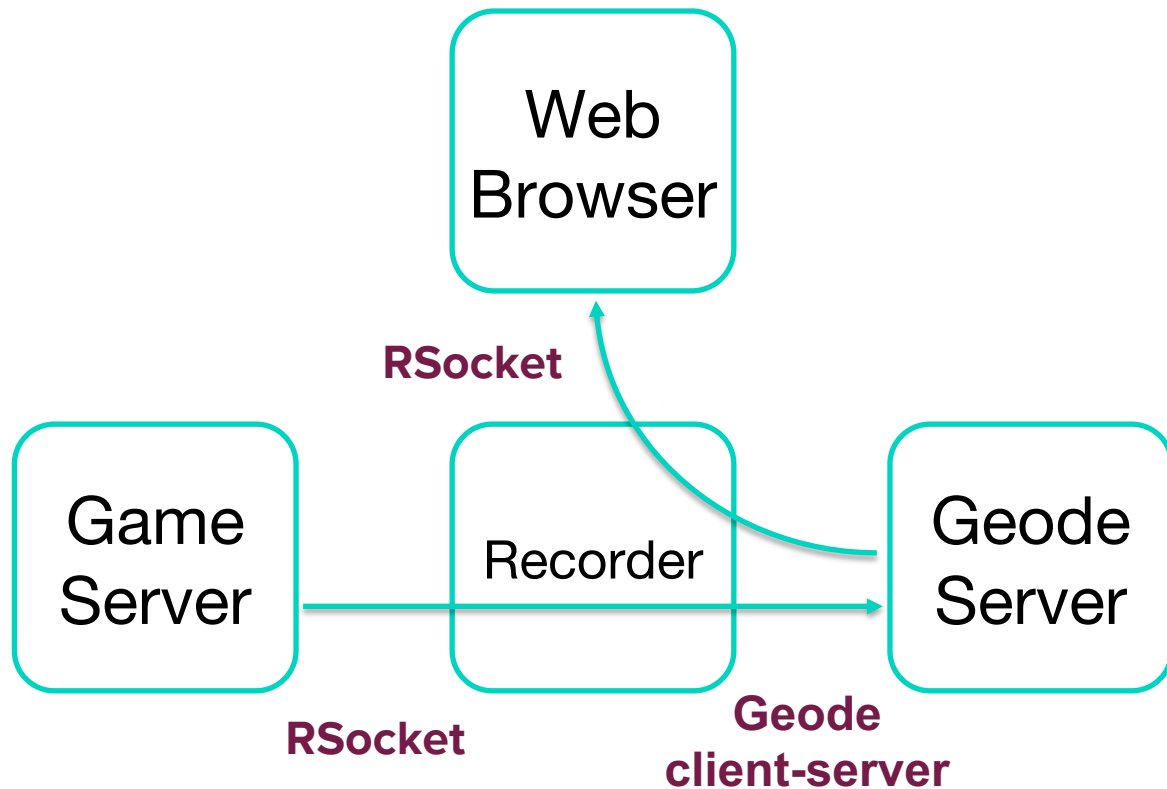
```
        // GeodeServerConfigurationAutoEvictionAndCQ.class,
```

```
        // "-Xmx200m", "-Xms200m",
```

```
        // While OpenJDK 12 defaults to G1GC now, Geode 1.9 doc says use CMS
```

```
        "-XX:+UseConcMarkSweepGC", "-XX:CMSInitiatingOccupancyFraction=60");
```

# Playback



# Hot Flux

```
82     private Publisher<Cell> getCellPublisherNewCellsOnly() {
83
84         return Flux.push(sink ->
85             {
86                 try {
87                     gemFireCache.getQueryService().newCq( queryString: "SELECT * FROM /Cells",
88                         createCqAttributes(sink)).execute();
89                 } catch (final CqException | RegionNotFoundException e) {
90                     sink.error(e);
91                 }
92             },
93             // BUFFERING == DANGER
94             FluxSink.OverflowStrategy.BUFFER);
95     }
96
97     @
98     private static CqListenerAdapter createCqListenerAdapter(final FluxSink<Cell> sink) {
99
100         return (cqEvent) -> {
101             if (cqEvent.getBaseOperation() == CREATE) {
102                 sink.next((Cell) cqEvent.getNewValue());
103             }
104         };
105     }
106
107     ...
```

# Endpoint

---

```
55 @RequestMapping("/rsocket/all-generations")
56 public Publisher<Cell> allGenerations(final Coordinates _ignored) {
57     return getCellPublisherNewCellsOnly();
58 }
```

# Results I

- Geode as reactive consumer (Subscriber)
  - How do active APIs fit with Geode's can Geode's load shedding policy be adapted to reactive back-pressure
    - Saw serial, parallel, `put`, `putAll`
    - `BlockHound + Schedulers.boundedElastic()`
  - Can Geode's load shedding policy be adapted to reactive back0pressure
    - App design, capacity planning, testing as usual
    - Back off on `LowMemoryException`

# Results II

---

- Geode as reactive producer (Publisher)
  - Again, how do reactive APIs *fit* with Geode's
    - Demonstrated a simple CQ -> hot Flux
  - Can Geode produce long data streams incrementally
    - Query results chunked cold/cold-ish Fluxes present some challenges
    - CQ ordering of initial results would be nice

Pivotal.

SpringOne Platform

by Pivotal.

# re:Cap Seoul 2019

2019년 12월 17일(화) • 한화 드림플러스 강남